

SRE Event Watchdog

Intelligent Observability & Event Watchdog

An **API-first** service that watches application logs in real time and uses a **hybrid statistical + ML + LLM** pipeline to catch incidents, explain their probable root cause, and fire alerts — **before a human notices**.

Fully local · one command · no API key required

Forward Deployed / GenAI Engineer challenge — built end-to-end with Claude Code

The problem: alert fatigue & late detection

On-call SREs drown in logs, and the tooling makes it worse:

- **Pure-threshold alerting** is noisy → engineers tune it out, real incidents slip through
- **Pure-ML** is a black box → a score with no "why," no action
- **Pure-LLM** is slow and expensive at log scale → can't sit in the hot path

Incidents get noticed late, and when they're noticed, the on-call still has to reverse-engineer *what broke* and *what to do*.

The bet: the right answer is **layered** — cheap detectors run on everything, the expensive LLM runs only on what already looks suspicious.

Architecture

```
flowchart LR
    GEN[Synthetic generator<br/>normal + injected incidents] -->|POST /api/logs| API
    CLIENT[Your services] -->|POST /api/logs| API
    subgraph APP[FastAPI app – single local process]
        API[Ingest] --> WIN[Windowing<br/>per-service buckets]
        WIN --> P{Hybrid detector}
        P -->|1 always-on| STATS[EWMA + z-score]
        P -->|2 multivariate| IF[Isolation Forest]
        P -->|3 gated, optional| LLM[Claude classifier]
        STATS --> SCORE[Hybrid score + severity]
        IF --> SCORE
        LLM --> SCORE
        SCORE --> DB[(SQLite)]
        SCORE --> ALERT[Alerter + cooldown]
        DB --> DASH[Chart.js dashboard]
    end
    ALERT -->|webhook| MOCK[Mock receiver<br/>Slack/PagerDuty stand-in]
    BROWSER[Browser] --> DASH
```

Logs → per-service time windows → 3-layer detection → persist → **alert + live dashboard.**

Why the hybrid design (the judgment)

Layer	Role	Why it's there	Cost
1. EWMA + z-score	Online statistical baseline	Instant signal, cold-start safe, interpretable. Baseline frozen during incidents so sustained outages keep flagging.	~free
2. Isolation Forest	Multivariate ML	Catches <i>contextual</i> anomalies a single-feature z-score misses. Flagged windows kept out of training.	cheap
3. Claude classifier	Root-cause triage	Turns a number into an explanation + recommended action. Gated behind layers 1–2.	gated

Cheap → expensive. The costly layer only runs on what's already suspicious — the way you'd actually run it in production.

The GenAI layer — production-shaped

The dimension the build is optimized to showcase:

- **Structured output via strict tool use** — a forced `report_root_cause` tool with an enum-constrained schema: `category`, `probable_root_cause`, `severity`, `recommended_action`, `confidence`. Machine-usable, not free text.
 - **Cost-gated** — only fires on windows the cheap detectors already flagged.
 - **Clean no-key fallback** — without `ANTHROPIC_API_KEY` the SDK isn't even imported; the app behaves identically on stats + ML. **It always runs.**
 - **Feeds back into alerting** — a `noise` verdict **suppresses the page.**
- | `claude-haiku-4-5` by default (fast/cheap for enrichment); model is configurable.

Live demo — one command

```
./scripts/run.ps1          # Windows    (./scripts/run.sh on macOS/Linux)
```

Brings up the API + dashboard (:8000) and the mock alert receiver (:8001).

Within **~1 minute**, autonomously:

1. The generator seeds healthy history → both detectors warm up in seconds
2. An incident is injected (error burst / latency regression / dependency outage)
3. Detectors trip → **hybrid** `stats+iforest` **anomaly, severity** `high`
4. An alert **POSTs to the mock receiver** (swap the URL for real Slack/PagerDuty)
5. The **dashboard lights up** — volume, error rate, latency p95, anomaly score

Trigger one on demand: `POST /api/demo/inject` — or click "Inject incident" in the UI.

Results

- **34 tests passing** — detector math, windowing, alerting + delivery, generator scenarios, dashboard contract, and the LLM layer (strict-tool-use, fallback, no-key path). Runs with **no key and no network**.
- **2-layer hybrid detection** confirmed live: injected incident → `iforest high score 0.96` → alert delivered to the mock receiver.
- **Measured tuning:** healthy false positives cut **~15% → ~1%** by measuring the z-score distribution (not guessing); residual ~1% left in honestly.
- **CI green** on Python 3.11 + 3.12 (ruff lint + pytest).
- **~10× test-suite speedup** (159s → 15s) from a thread-local SQLite pool.

Vibe-coding process — engineering judgment

Built **end-to-end with Claude Code under a "no manual edits" rule**; full prompt-by-prompt audit trail in `prompts.md`, spec in `SPEC.md`. Four real debugging calls show the difference between *generating code* and *engineering a system*:

- **EWMA `alpha` tuning** — a test showed slow convergence; raised 0.3→0.4 by reasoning, **not** by gaming the test.
- **Isolation Forest dead end** — added a `StandardScaler`, proved it a no-op (IF is scale-invariant), **reverted it**, found the real cause (constant test features).
- **False-positive rate** — *measured* the z-distribution, raised the threshold, and **left the honest ~1% residual** rather than fake a zero.
- **httpx 0.28 `ASGITransport`** — root-caused an async-only SDK change; fixed in the test, not production.

Production hardening — what's next

A focused MVP. To run it for real:

- **Scale:** queue in front of ingestion; move windows off SQLite to Postgres/Timescale.
- **Detector lifecycle:** persist/version models, scheduled retraining, drift monitoring, backtesting.
- **LLM layer:** response caching + token budgets, prompt/version pinning, an eval harness, human feedback to tune categories.
- **Alerting:** real Slack/PagerDuty/Opsgenie, dedup/grouping, escalation, ack tracking.
- **Reliability & security:** auth + rate limiting on ingest, metrics for the watchdog itself, readiness probes.
- **Delivery:** Dockerfile + compose, Helm chart. *(CI already in place.)*

Thank you — github.com/mallimatla/sre-event-watchdog